

Meeting Discussion 23rd May 2021

Lunar Lander Pseudo Code

1. Initialize Policy parameters θ and State Value function parameters W
2. For each episode:

1. Initialize State S_0 (first state of the episode)
2. While S is not terminal:
 1. Sample action a_t , based on actor's policy μ_θ
$$a \sim \pi(S, \cdot, \theta)$$
 2. Receive reward R_{t+1} and update to next state (S_t to S_{t+1}).
 3. Save the action and the Value of the state into a list.

$$S_a.insert(\ln\mu_\theta(s_t, a_t), V_w(s_t))$$

where,

$\ln\mu_\theta(s_t, a_t)$ = Log probability of selecting an action at time t

$V_w(s_t)$ = Value of the state at time t

S_a = List of saved actions and Value of the state

4. Add the rewards to the episode's total rewards

$$E_r = E_r + R_{t+1}$$

where,

E_r = Episode Reward,

3. Update Average Reward:

$$avg_r = (E_r + avg_r)/2$$

where,

avg_r = Average Reward,

4. Update the Loss functions and the parameters of critic and actor:

1. Compute all n-step returns of the episode

$$G_t = \sum_{k=0}^{n-t} \gamma^k R_{k+t+1}$$

Where,

G_t = t step total return

R = Reward

2. Compute the total Policy Loss, i.e.

$$\delta_{policy\ loss} = \sum_{t=0}^{t=n} (-\ln \mu_{\theta}(s_t, a_t)(G_t - V_w(s_t)))$$

3. Compute the total Value loss, i.e.

$$\delta_{value\ loss} = \sum_{t=0}^{t=n} s\ smooth_{L_1}(V_w(s_t), G_t)$$

4. Compute total loss, i.e.

$$total_{loss} = \delta_{policy\ loss} + \delta_{value\ loss}$$

5. Update the parameters based on $total_{loss}$

6. Hence updating the weights θ and w of the networks.

```
# saved_actions = list of (log_prob of actions selectected, and Value of state)
# returns = list of t-step return

for (log_prob, value), R in zip(saved_actions, returns):
    # Advantage function calculated
    advantage = R - value.item()
    # Actor's loss appended to sum later
    ActorCritic_losses.append((-log_prob) * advantage)
    # Critic's loss appended to sum later
    value_losses.append(F.smooth_l1_loss(value, torch.tensor([R])))

# Clears the previous episodes gradients
optimizer.zero_grad()
# Total loss calculated which is distributed to the respective parameters
# by pytorch itself.
loss = torch.stack(ActorCritic_losses).sum() + torch.stack(value_losses).sum()
# Perform backward propagation
loss.backward()
# Perform gradient descent (both parameter's update)
optimizer.step()
```